

GCS Snowflake Object Inventory — Structure and Description

Database: CUSTOMER_SUPPORT_BOT_LOGS | Schema: CHAT_DATA

Object name	Type / status	Structure	Description and update behavior
DBD_OWNER_DATA	Reference table / existing	Key: ID NAME, REGION	Maps Salesforce support-owner names to operating regions. Used to define dashboard owners, region filters, customer-vs-support comment authors, and notification scope.
DBD_SENTIMENT_DATA	Current-state table / existing	Key: CASENUMBER ID, SENTIMENT, CREATED_AT	Stores the final four-category sentiment used by the dashboard. The hourly OpenAI pipeline inserts new cases and updates SENTIMENT for existing cases; CREATED_AT remains the original insert time.
SENTIMENT_ANALYSIS_RESULTS_72H	Reference/history table / existing	ID, RUN_ID, RUN_TIMESTAMP, CASE_ID, CASE_NUMBER, CASE_SUBJECT, CASE_STATUS, CASE_PRIORITY, OWNER_NAME, ACCOUNT_NAME, CASE_CREATED_DATE, COMMENT_COUNT, SENTIMENT, SENTIMENT_SCORE, SENTIMENT_SCORE_100, LATEST_CUSTOMER_COMMENT, SUMMARY, RECOMMENDED_ACTION, COMMENTS_ANALYZED_JSON, CREATED_AT	External 72-hour sentiment source. The project selects exactly the newest row per CASE_NUMBER using CREATED_AT descending (ID descending tie-breaker), maps its label to Positive/Neutral/Negative/Critical, and compares it with OpenAI. OpenAI remains authoritative on mismatch.
OPENAI_LLM_USAGE	Usage audit table / existing	Key: USAGE_ID CREATED_AT, AGENT_NAME, OPERATION, STATUS, MODEL, RESPONSE_ID, LATENCY_MS, INPUT_TOKENS, CACHED_INPUT_TOKENS, OUTPUT_TOKENS, REASONING_TOKENS, TOTAL_TOKENS, ESTIMATED_COST_USD, PRICING_TIER, PRICE_SOURCE, RAW_USAGE_JSON, METADATA_JSON, ERROR_MESSAGE	Captures OpenAI request status, latency, token consumption, estimated cost, response identifiers, metadata, raw usage and failures for batch and fallback sentiment calls.
DBD_CASE_AUDIT_HISTORY	Historical audit table / existing	Key: AUDIT_ID CASE_NUMBER, SNAPSHOT_TIMESTAMP, CHANGE_TYPE, CHANGED_COLUMNS, OLD_STATE, NEW_STATE, DATA_HASH	Stores dashboard case-state history and changed-column details. The dashboard compares hashes against the latest stored snapshot and inserts new audit rows; configured with short retention and clustering by timestamp/case.
DBD_SLA_BREACH_IMPACT	Historical impact table / existing	Key: SLA_IMPACT_ID RECORD_TYPE, SNAPSHOT_DATE, SNAPSHOT_TIMESTAMP, CASE_NUMBER, CUSTOMER_NAME, CASE_OWNER, SEVERITY, SUPPORT_LEVEL, SUPPORT_TIER, SALESFORCE_STATUS, SLA_STATUS, SLA_DEADLINE, LAST_CUSTOMER_COMMENT, SLA_VARIANCE_MINS, BREACH_MONTH, IS_HEAL_DESK, SUBJECT, IMPACT_STATUS, IMPACT_REASON, PREVIOUS_SLA_STATUS, CURRENT_SLA_STATUS, CREATED_AT	Maintains current and historical SLA-breach impact records for reporting. Updated by the SLA synchronization process with deduplication and status-transition context.
SF_CASE_ANALYST	Reference table / existing	Key: SALESFORCE_USER_ID ANALYST_NAME, EMAIL, IS_ACTIVE, MANAGER_NAME, LAST_UPDATED	Maps case-owner names to active Slack recipient email addresses. Exact analyst-name matching is required; the newest active mapping is selected using LAST_UPDATED.
DBD_GCS_NOTIFICATION_RUNS	Scheduler control table / existing	Key: RUN_ID STARTED_AT, COMPLETED_AT, STATUS, STATS, ERROR_MESSAGE	Claims deterministic shift-digest runs and prevents duplicate sends for the same scheduled region/time. Records completion statistics or failure details.
DBD_GCS_CASE_NOTIFICATION_AUDIT	Historical event table / existing	Key: EVENT_ID; Grouping: DELIVERY_ID EVENT_AT, USE_CASE, NOTIFICATION_TYPE, CASE_NUMBER, ANALYST_NAME, ANALYST_EMAIL, REGION, CUSTOMER_NAME, PREVIOUS_OWNER, CURRENT_OWNER, PREVIOUS_PRIORITY, CURRENT_PRIORITY, PREVIOUS_STATUS, CURRENT_STATUS, PREVIOUS_ESCALATED, CURRENT_ESCALATED, SLACK_USER_ID, SLACK_MESSAGE_TS	Stores case events only after successful Slack delivery for use cases 2 and 3. Shift digests are excluded. Multiple cases in one consolidated Slack message share DELIVERY_ID.
DBD_GCS_CASE_NOTIFICATION_SUMMARY	Reporting view / project DDL (not present)	PERIOD_TYPE, PERIOD_START, ANALYST_NAME, USE_CASE, NOTIFICATION_TYPE, DISTINCT_CASE_COUNT, SLACK_MESSAGE_COUNT, EVENT_COUNT	Planned management view over notification audit history, producing daily, weekly and monthly IST counts by analyst, use case and notification type.
DBD_GCS_CASE_PRIORITY_SNAPSHOT	Current-state table / existing	Key: CASE_NUMBER CASE_ID, REGION, CUSTOMER_NAME, CASE_OWNER, SUPPORT_LEVEL, SEVERITY, STATUS, ESCALATED, SENTIMENT, LAST_COMMENT_BY, SLA_RESPONSE_TIME, SLA_MINUTES, SEVONE, PRIORITY, CASE_SCORE, SUBJECT, IS_HEAL_DESK, SNAPSHOT_AT, CREATED_AT, UPDATED_AT	Authoritative current Case Priority Index dataset for downstream systems. Every successful background refresh globally re-ranks cases, upserts present rows and removes cases no longer returned by the dashboard source.

2. Documentation and Operational Usage

2.1 Usage for Engineers and Analysts Use the Case Priority Index as the working queue for active support cases. Start with the lowest PRIORITY number, then review severity, escalation, sentiment, latest-comment source and SLA deadline before acting. The browser keeps its current snapshot until manual Refresh, while the backend refreshes every 60 seconds. DBD_GCS_CASE_PRIORITY_SNAPSHOT is the machine-readable current state for integrations. Sentiment is calculated from only the latest published customer comment; cases without a customer comment are Neutral. OpenAI remains authoritative when the mapped 72-hour sentiment differs.	2.2 Usage for Queue Management Use REGION, CASE_OWNER, STATUS and PRIORITY to balance workload and identify ownership gaps. New owner assignments and transitions from any status to Assigned generate the use-case-two Slack queue. The new/follow-up case and any rows whose rank changed are bold. SF_CASE_ANALYST must contain an active exact-name email mapping for delivery. DBD_GCS_CASE_NOTIFICATION_AUDIT and its summary view support analyst-level daily, weekly and monthly assignment/follow-up counts.
2.3 Usage for SLA Cleanups Sort and filter using SLA_RESPONSE_TIME and SLA_MINUTES to identify overdue or near-due work. Use DBD_SLA_BREACH_IMPACT for historical breach state, impact reason and SLA transitions; use DBD_GCS_CASE_PRIORITY_SNAPSHOT for the current actionable queue. Clear overdue customer-facing actions first, then validate status, latest comment and ownership in Salesforce. Slack is an alerting aid and does not replace Salesforce as the case system of record.	2.4 Understanding Prioritization for Leadership Priority is a rank, not a manually entered severity. The dashboard calculates Case Score from escalation, SevOne, low-priority record type, severity, support tier, customer sentiment and SLA overdue state/duration. Current precedence is: Escalated = 30; otherwise low-priority record types = 0.25; otherwise SevOne = 35; remaining cases use the scoring matrix and overdue bonus. Cases are globally sorted by Case Score descending and SLA minutes ascending, then numbered Priority 1, 2, 3, and so on. A lower priority number means earlier attention.

Slack Alert Scenarios

Scenario	Trigger	Message behavior	Tracking
1 — Shift-start priority digest	Weekday shift start: APAC 6:00 AM, EMEA 2:00 PM, NA East 6:00 PM, NA West 9:00 PM IST.	One complete priority table per eligible owner. Current recipient scope is NA East. Duplicate runs are prevented by DBD_GCS_NOTIFICATION_RUNS.	Not written to the case-notification audit table.
2 — New case / follow-up response	A case is assigned/reassigned to an owner, or transitions from any other status to Assigned.	One consolidated owner queue. The new/follow-up case and rows whose priorities changed are bold.	Recorded as USE_CASE_2 with NEW_ASSIGNMENT or FOLLOW_UP_RESPONSE after successful Slack delivery.
3 — Priority elevation / escalation	A case priority changes or its escalation state changes between background snapshots.	One consolidated owner queue. Only rows moving to a better rank and newly escalated rows are bold; downward shifts remain normal.	Recorded as USE_CASE_3 with priority and escalation event types after successful Slack delivery.

Operational note: the first background snapshot after startup establishes a baseline and does not send change alerts. Subsequent comparisons run approximately every 60 seconds.