GCS Snowflake Object Inventory — Structure and Description

Database: CUSTOMER_SUPPORT_BOT_LOGS | Schema: CHAT_DATA

Object name	Type / status	Structure	Description and update behavior
DBD_OWNER_DATA	Reference table / existing	**Key:** ID NAME, REGION	Maps Salesforce support-owner names to operating regions. Used to define dashboard owners, region filters, customer-vs-support comment authors, and notification scope.
DBD_SENTIMENT_DATA	Current-state table / existing	**Key:** CASENUMBER ID, SENTIMENT, CREATED_AT	Stores the final four-category sentiment used by the dashboard. The hourly OpenAI pipeline inserts new cases and updates SENTIMENT for existing cases; CREATED_AT remains the original insert time.
SENTIMENT_ANALYSIS_RESULTS_72H	Reference/history table / existing	ID, RUN_ID, RUN_TIMESTAMP, CASE_ID, CASE_NUMBER, CASE_SUBJECT, CASE_STATUS, CASE_PRIORITY, OWNER_NAME, ACCOUNT_NAME, CASE_CREATED_DATE, COMMENT_COUNT, SENTIMENT, SENTIMENT_SCORE, SENTIMENT_SCORE_100, LATEST_CUSTOMER_COMMENT, SUMMARY, RECOMMENDED_ACTION, COMMENTS_ANALYZED_JSON, CREATED_AT	External 72-hour sentiment source. The project selects exactly the newest row per CASE_NUMBER using CREATED_AT descending (ID descending tie-breaker), maps its label to Positive/Neutral/Negative/Critical, and compares it with OpenAI. SENTIMENT_ANALYSIS_RESULTS_72H is authoritative on mismatch.
OPENAI_LLM_USAGE	Usage audit table / existing	**Key:** USAGE_ID CREATED_AT, AGENT_NAME, OPERATION, STATUS, MODEL, RESPONSE_ID, LATENCY_MS, INPUT_TOKENS, CACHED_INPUT_TOKENS, OUTPUT_TOKENS, REASONING_TOKENS, TOTAL_TOKENS, ESTIMATED_COST_USD, PRICING_TIER, PRICE_SOURCE, RAW_USAGE_JSON, METADATA_JSON, ERROR_MESSAGE	Captures OpenAI request status, latency, token consumption, estimated cost, response identifiers, metadata, raw usage and failures for batch and fallback sentiment calls.
DBD_CASE_AUDIT_HISTORY	Historical audit table / existing	**Key:** AUDIT_ID CASE_NUMBER, SNAPSHOT_TIMESTAMP, CHANGE_TYPE, CHANGED_COLUMNS, OLD_STATE, NEW_STATE, DATA_HASH	Stores dashboard case-state history and changed-column details. The dashboard compares hashes against the latest stored snapshot and inserts new audit rows; configured with short retention and clustering by timestamp/case.
DBD_SLA_BREACH_IMPACT	Historical impact table / existing	**Key:** SLA_IMPACT_ID RECORD_TYPE, SNAPSHOT_DATE, SNAPSHOT_TIMESTAMP, CASE_NUMBER, CUSTOMER_NAME, CASE_OWNER, SEVERITY, SUPPORT_LEVEL, SUPPORT_TIER, SALESFORCE_STATUS, SLA_STATUS, SLA_DEADLINE, LAST_CUSTOMER_COMMENT, SLA_VARIANCE_MINS, BREACH_MONTH, IS_HEAL_DESK, SUBJECT, IMPACT_STATUS, IMPACT_REASON, PREVIOUS_SLA_STATUS, CURRENT_SLA_STATUS, CREATED_AT	Maintains current and historical SLA-breach impact records for reporting. Updated by the SLA synchronization process with deduplication and status-transition context.
SF_CASE_ANALYST	Reference table / existing	**Key:** SALESFORCE_USER_ID ANALYST_NAME, EMAIL, IS_ACTIVE, MANAGER_NAME, LAST_UPDATED	Maps case-owner names to active Slack recipient email addresses. Exact analyst-name matching is required; the newest active mapping is selected using LAST_UPDATED.
DBD_GCS_NOTIFICATION_RUNS	Scheduler control table / existing	**Key:** RUN_ID STARTED_AT, COMPLETED_AT, STATUS, STATS, ERROR_MESSAGE	Claims deterministic shift-digest runs and prevents duplicate sends for the same scheduled region/time. Records completion statistics or failure details.
DBD_GCS_CASE_NOTIFICATION_AUDIT	Historical event table / existing	**Key:** EVENT_ID; **Grouping:** DELIVERY_ID EVENT_AT, USE_CASE, NOTIFICATION_TYPE, CASE_NUMBER, ANALYST_NAME, ANALYST_EMAIL, REGION, CUSTOMER_NAME, PREVIOUS_OWNER, CURRENT_OWNER, PREVIOUS_PRIORITY, CURRENT_PRIORITY, PREVIOUS_STATUS, CURRENT_STATUS, PREVIOUS_ESCALATED, CURRENT_ESCALATED, SLACK_USER_ID, SLACK_MESSAGE_TS	Stores case events only after successful Slack delivery for use cases 2 and 3. Shift digests are excluded. Multiple cases in one consolidated Slack message share DELIVERY_ID.
DBD_GCS_CASE_NOTIFICATION_SUMMARY	Reporting view / existing	PERIOD_TYPE, PERIOD_START, ANALYST_NAME, USE_CASE, NOTIFICATION_TYPE, DISTINCT_CASE_COUNT, SLACK_MESSAGE_COUNT, EVENT_COUNT	Planned management view over notification audit history, producing daily, weekly and monthly IST counts by analyst, use case and notification type.
DBD_GCS_CASE_PRIORITY_SNAPSHOT	Current-state table / existing	**Key:** CASE_NUMBER CASE_ID, REGION, CUSTOMER_NAME, CASE_OWNER, SUPPORT_LEVEL, SEVERITY, STATUS, ESCALATED, SENTIMENT, LAST_COMMENT_BY, SLA_RESPONSE_TIME, SLA_MINUTES, SEVONE, PRIORITY, CASE_SCORE, SUBJECT, IS_HEAL_DESK, SNAPSHOT_AT, CREATED_AT, UPDATED_AT	Authoritative current Case Priority Index dataset for downstream systems. Every successful background refresh globally re-ranks cases, upserts present rows and removes cases no longer returned by the dashboard source.

2. User Guide

2.1 Usage for Engineers / Analysts Open the Case Priority Index to see active cases in recommended working order. Cases with Priority 1 should normally be reviewed first. Before acting, review severity, escalation, sentiment, SLA response time, latest comment and owner. Confirm important details in Salesforce. Use the region, owner, status, support-level and Heal Desk filters to focus the list. Select Refresh whenever you need the latest data. Sentiment reflects the latest customer comment. When no customer comment is available, the case is displayed as Neutral. The 72-hour sentiment result is authoritative and overrides the OpenAI classification when the two results differ.	**2.2 Usage for Queue Management** Filter by region, owner and status to compare workload and find unassigned or overloaded queues. Review top-priority cases first. Before reassigning work, check severity, sentiment, SLA position and the owner's active workload. In Slack queue alerts, bold rows identify the new or follow-up case and other cases whose queue position changed. After making an ownership change, confirm the new owner and status in Salesforce.
2.3 Usage for SLA Cleanups Use SLA Response Time to find overdue and near-due cases. Review red overdue cases immediately and plan due-soon work next. Open the case in Salesforce. Read the latest customer comment, confirm ownership and decide whether a response is required. Work overdue customer-facing actions first, followed by the closest upcoming deadlines. Keep the Salesforce status and owner current. Use dashboard and Slack alerts as guidance; Salesforce remains the case system of record.	**2.4 Understanding Prioritization for Leadership** Priority is the recommended working order, not Salesforce severity. A lower number means earlier review. The ordering considers escalation, SevOne, severity, support level, customer sentiment and whether the SLA is overdue. Critical and Negative sentiment raise attention, while overdue duration helps order cases with similar business importance. Always apply support judgment and confirm the latest Salesforce information before acting.

Slack Alerts and Recommended Actions

Scenario	When it appears	What you will see	What you should do
1 — Shift-start priority digest	At the scheduled start of each weekday support shift.	A complete prioritized list of active cases for each eligible owner.	Review the top cases and plan the shift workload.
2 — New case / follow-up response	A case is assigned, reassigned or moved to Assigned status.	The owner's updated queue; new, follow-up and changed-priority rows are bold.	Review the bold case and respond if customer action is due.
3 — Priority elevation / escalation	A case moves in priority or becomes escalated.	The updated queue; higher-priority moves and new escalations are bold.	Review the bold cases promptly and update the action plan in Salesforce.

Tip: Refresh the dashboard before making queue decisions, and verify important case details in Salesforce before taking action.